

Copy-on-Write Fork for xv6

Implement copy-on write (hard)

The general idea is to mark all pagetables as not writable and with bits for recognizing copy-on-write (COW) set to 1. Whenever a page fault occurs, the trap handling function should be able to identify this cause and update the pagetable.

This allows the copy process of fork to happen only when the process needs to write corresponding memory, which would save a lot of time when forking a very large process.

To implement this functionality, the instructions on the [lab website](#) give a list of functions that should be modified:

- `uvmcopy` : remove the memory allocation and map the parent to the child
- `usertrap` : where we handle the page fault, the core part of this task
- `kalloc` and `kfree` : manage the page's reference count
- `kinit` : initialize the page's reference count
- `copyout` : reuse the scheme in `usertrap` to handle COW pages

To start with, we will update the `uvmcopy` function. We need to remove the memory allocation part of the original code. An initial attempt at this is shown below:

```
int
uvmcopy(pagetable_t old, pagetable_t new, uint64 sz)
{
    pte_t *pte;
    uint64 pa, i;
    uint flags;

    for(i = 0; i < sz; i += PGSIZE){
        // The reason we use walk here is to validate the PTE flags.
        // So that we can make sure the PTE is valid and needs to be copied.
        if((pte = walk(old, i, 0)) == 0)
            panic("uvmcopy: pte should exist");
        if((*pte & PTE_V) == 0)
            panic("uvmcopy: page not present");
    } // examine the old pagetable
    for (i = 0; i < 512; i++) {
        *pte = &old[i];          // get pointer to old page table entry
        *pte = *pte & ~PTE_W;    // clear write bit
        *pte = *pte | PTE_COW;   // set COW bit
        new[i] = *pte;           // copy page table entry directly
        // increase the page reference count
        page_reference_count[PGREF_CNT(pa)]++;
    }
    return 0;
}
```

We can remove the rest of the original code that uses the `mappages` function or allocates memory. This is because we can update the PTE manually later in the given `for` loop.

In the `copyout` function, we can insert a branch to check for `COW` pages marked by `PTE_COW`. The original code in

this function uses `walkaddr` instead of `walk` to avoid potential vulnerabilities when writing to the `pte`. However, in this case we need to update the `pte` according to the `pte` flags, so we can replace the `walkaddr` function with the `walk` function and obtain a returned value of type `pte_t *`. We can then update the page reference count based on the count of the specific physical address. There is no need to use the `mappages` function here, as we can directly modify the value of a specific `pte` using its pointer.

The `page_reference_count` array should be declared using the `extern` keyword to make the global variable accessible in this file. It is defined in `kalloc.c`.

```
int
copyout(pagetable_t pagetable, uint64 dstva, char *src, uint64 len)
{
    uint64 n, va0, pa0;
    pte_t *pte;

    while(len > 0){
        va0 = PGROUNDDOWN(dstva);
        pte = walk(pagetable, va0, 0);
        pa0 = PTE2PA(*pte);
        if(pa0 == 0)
            return -1;
        n = PGSIZE - (dstva - va0);
        if(n > len)
            n = len;
        if (*pte & PTE_COW) {
            // if the page is COW, we need to copy the page first
            // before we can write to it
            uint flags = PTE_FLAGS(*pte);
            char *page;
            if((page = kalloc()) == 0)
                panic("copyout(): out of memory");
            memmove(page, (char*)pa0, PGSIZE);
            *pte = PA2PTE((uint64)page) | (flags & ~PTE_COW) | PTE_W;
            // decrease the page reference count
            if (page_reference_count[PGREF_CNT(pa0)] <= 1) {
                page_reference_count[PGREF_CNT(pa0)] = 0;
                pa0 = (uint64) page;
            } else {
                page_reference_count[PGREF_CNT(pa0)]--;
            }
        }
        // original code here
        memmove((void *)(pa0 + (dstva - va0)), src, n);
        ...
    }
    return 0;
}
```

In `kalloc.c`, there are several slight modifications that should be applied. First, define the `page_reference_count` array with a size of `(PHYSTOP-KERNBASE) / PGSIZE`, as suggested by the `cowtest.c` file shown below:

```
void
simpletest()
{
    uint64 phys_size = PHYSTOP - KERNBASE;
    // other part of code
    ...
}
```

This is slightly different from what is mentioned on the [lab website](#):

"You'll have to work out a scheme for how to index the array and how to choose its size. For example, you could index the array with the page's physical address divided by 4096, and give the array a number of elements equal to highest physical address of any page placed on the free list by kinit() in kalloc.c."

However, the currently implemented approach theoretically saves more memory.

The next modification to be made in `kalloc.c` is applied to the `kalloc` function. We should set the reference count of each newly allocated memory block to 1.

```
void *
kalloc(void)
{
    struct run *r;
    // other part of code
    ...
    if(r){
        memset((char*)r, 5, PGSIZE); // fill with junk
        page_reference_count[PGREF_CNT((uint64)r)] = 1; // inserted here
    }

    return (void*)r;
}
```

Later on, we need to focus on the very important part of this task: handling the page fault in the `usertrap` function. We use the `r_scause()` function to retrieve the value from the RISC-V `scause` register, which is used to store the type of fault. We can refer to the RISC-V manual to find the appropriate `scause` value for a page fault, as shown [here](#) or in the table below:

Interrupt	Exception Code	Description
0	12	Instruction page fault
0	13	Load page fault
0	14	Reserved
0	15	Store/AMO page fault

From the descriptions, we can understand that the page fault `scause` value for `cow` handling should be `12`, `13`, or `15` (actually, it should only be `15` because `cow` pages should only be handled during write operations, which correspond to `store` page faults, but it won't influence the program's behavior in this lab).

```

if(r_scause() == 8){
    // handling syscall
    ...
    syscall();
} else if(r_scause() == 12 || r_scause() == 13 || r_scause() == 15){
    // page fault
    uint64 va = r_stval(); // store the faulting address
    // free previous page and add the newly allocated page to the pagetable
    uint64 *pte = walk(p->pagetable, va, 0);
    if (pte == 0) {
        panic("usertrap(): pte should exist\n");
    }
    if ((*pte & PTE_V) == 0) {
        panic("usertrap(): page should exist\n");
    }
    if ((*pte & PTE_COW)) {
        uint64 pa = PTE2PA(*pte);
        uint flags = PTE_FLAGS(*pte);
        char* page; // allocate a new page
        if ((page = kalloc()) == 0) {
            p->killed = 1;
            goto trap_err;
        }
        memmove(page, (char*)pa, PGSIZE); // copy the content of the old page to the new page
        *pte = PA2PTE((uint64)page) | (flags & ~PTE_COW) | PTE_W;

        // free the old page if the last reference to it is removed
        if (page_reference_count[PGREF_CNT(pa)] <= 1) {
            page_reference_count[PGREF_CNT(pa)] = 0;
            pa = (uint64)page;
        } else {
            page_reference_count[PGREF_CNT(pa)]--;
        }
    }
} else if((which_dev = devintr()) != 0){
    // ok
} else {
    printf("usertrap(): unexpected scause %p pid=%d\n", r_scause(), p->pid);
    printf("          sepc=%p stval=%p\n", r_sepc(), r_stval());
    p->killed = 1;
}
trap_err:
...
// the rest of the function

```

The idea applied here is similar to that used in the `copyout` function, as instructed. We still use the pointer to the `pte` so that we don't need to call the `mappages` function to `walk` through the `va` again.

However, something went wrong in the `usertests` even though the `cowtest` passed. It resulted in a kernel trap when running the `usertests` `copyout`. I tried several approaches here, including moving the page reference count into the `kfree` function as many bloggers do, and handling the MAXVA problem in the page fault handling section.

The updated code is explained below, although it still only passes the `cowtest` and gets stuck even at the first `usertests`, which makes things worse, but I forgot to commit my previous code here, only the commented code line remains.

First, I rewrote my `uvmcopy` function as many bloggers do - I just removed the allocation part but kept the `mappages` function here. This is because it is okay to use `mappages` here to ensure the correctness of the code before passing the `usertests`. That is, I made as few changes to the original code as possible for now.

```

int
uvmcopy(pagetable_t old, pagetable_t new, uint64 sz)
{
    pte_t *pte;
    uint64 pa, i;
    uint flags;

    for(i = 0; i < sz; i += PGSIZE){
        // The reason we use walk here is to validate the PTE flags.
        // So that we can make sure the PTE is valid and needs to be copied.
        if((pte = walk(old, i, 0)) == 0)
            panic("uvmcopy: pte should exist");
        if((*pte & PTE_V) == 0)
            panic("uvmcopy: page not present");
        pa = PTE2PA(*pte);
        flags = (PTE_FLAGS(*pte) & (~PTE_W)) | PTE_COW;

        *pte = PA2PTE(pa) | flags;
        if(mappages(new, i, PGSIZE, pa, flags) != 0)
            goto err;
        if(((uint64)pa % PGSIZE) != 0 || (char*)pa < end || (uint64)pa >= PHYSTOP)
            goto err;
        page_reference_count[PGREF_CNT(pa)]++;
    }
    return 0;

err:
    uvmunmap(new, 0, i / PGSIZE, 1);
    return -1;
}

```

Next, we updated the `copyout` function and moved the page reference count behavior into the `kfree` function. This allows us to handle the reference count in an implicit way. In the branch of the `if` statement that handles `cow` pages, we replaced the decrease page reference count part with a single `kfree` statement.

```

int
copyout(pagetable_t pagetable, uint64 dstva, char *src, uint64 len)
{
    uint64 n, va0, pa0;
    pte_t *pte;
    while(len > 0){
        ...
        if (*pte & PTE_COW) {
            // other part of handling
            ...
            // decrease the page reference count
            // if (page_reference_count[PGREF_CNT(pa0)] <= 1) {
            //     page_reference_count[PGREF_CNT(pa0)] = 0;
            kfree((void*)pa0);
            pa0 = (uint64) page;
            // } else {
            //     page_reference_count[PGREF_CNT(pa0)]--;
            // }
        }
        ...
    }
    ...
}

```

Now, start modifying the `kalloc.c` functions. First is absolutely the `kfree` function where we should handling the

reference count, the logic will not change compared to the code implementation in `usertrap`.

```
void
kfree(void *pa)
{
    struct run *r;

    if(((uint64)pa % PGSIZE) != 0 || (char*)pa < end || (uint64)pa >= PHYSTOP)
        panic("kfree");

    // inserted code here
    if (page_reference_count[PGREF_CNT((uint64)pa)] > 1) {
        page_reference_count[PGREF_CNT((uint64)pa)]--;
        return;
    } else {
        page_reference_count[PGREF_CNT((uint64)pa)] = 0;
    }

    // Fill with junk to catch dangling refs.
    ...
}
```

We also modified the `freerange` function. In the previous version, we didn't need to initialize the `page_reference_count` array with `1` because we didn't use `kfree` to handle the page reference count. Instead, we did it manually in the `usertrap` and `copyout` functions. The initial value of global variables is `0` by default.

```
void
freerange(void *pa_start, void *pa_end)
{
    char *p;
    p = (char*)PGROUNDUP((uint64)pa_start);
    for(; p + PGSIZE <= (char*)pa_end; p += PGSIZE){
        page_reference_count[PGREF_CNT((uint64)p)] = 1; // inserted code here
        kfree(p);
    }
}
```

Finally, we modified the code in `usertrap` (after updating the `copyout` function). We manually managed the `page_reference_count` array in a similar way as we did in `copyout`.

However, the program still failed on the first `usertest`, which is unacceptable for us. After looking into the issue, we found that the program was stuck at `*(char*)a = 99;` in the child process during the `MAXVAp1us` test. The for loop was still on the first round, and it was easy to see that the address `a` was out of range because it was initially set to `MAXVA`. This issue should have been handled in `usertrap` when the virtual address was invalid. We found that the line `if(va >= MAXVA || (va <= PGROUNDDOWN(p->trapframe->sp) && va >= PGROUNDDOWN(p->trapframe->sp)-PGSIZE))` was causing an infinite loop by directly going to the `trap_err` label without updating the `killed` field of the process to `1`.

When we fixed this bug, the `usertests` were able to run until the `copyout` function, but we still encountered a kernel trap. Unfortunately, due to time constraints, we were not able to fully resolve this issue before writing reports of other labs.

Make grade

```
make[1]: Leaving directory '/home/lydia/projects/xv6-labs-2021'
```

```

== Test running cowtest ==
$ make qemu-gdb
(11.1s)
== Test    simple ==
    simple: OK
== Test    three ==
    three: OK
== Test    file ==
    file: OK
== Test usertests ==
$ make qemu-gdb
Timeout! (300.2s)
    ...
    test copyin: OK
    test copyout: scause 0x000000000000000d
    sepc=0x00000000080000c20 stval=0x0000000000000000
    panic: kerneltrap
    qemu-system-riscv64: terminating on signal 15 from pid 3859363 (make)
MISSING '^ALL TESTS PASSED$'
QEMU output saved to xv6.out.usertests
== Test    usertests: copyin ==
    usertests: copyin: FAIL
    Parent failed: test_usertests
== Test    usertests: copyout ==
    usertests: copyout: FAIL
    Parent failed: test_usertests
== Test    usertests: all tests ==
    usertests: all tests: FAIL
    Parent failed: test_usertests
== Test time ==
time: FAIL
    Cannot read time.txt
Score: 80/110
make: *** [Makefile:336: grade] Error 1
lydia@ubuntu-22-hp-040f1b4d: ~/projects/xv6-labs-2021$

```

Reference

1. [xyfJASON's GitHub](#)
2. [NebulorDang's GitHub](#)
3. [巴勃罗·捏捏达's CSDN](#)
4. [MIT 6.S081 Lab: Copy-on-Write Fork for xv6](#)