

Multithreading

Uthread: switching between threads (moderate)

In this task, our goal is to create a user function that can support multithreading, similar to how library developers have implemented it in various programming languages. It can be challenging to handle the low-level implementation of multithreading, but for this lab, the task is relatively straightforward.

To begin, we should read through the code of several `xv6` functions, as outlined in the `xv6` book. These functions include `yield`, `sleep`, `sched`, `scheduler`, and the assembly code `swtch`. These resources will give us a good understanding of how scheduling works within the C code, which will be helpful as we start implementing this task.

The instructions provided may not be very clear, but the task itself is fairly simple to accomplish by just reading through the `uthread.c` code example. Following the TODOs, our first step will be to complete the definition of `struct thread` using the following code:

```
struct thread {
    char    stack[STACK_SIZE]; /* the thread's stack */
    int     state;              /* FREE, RUNNING, RUNNABLE */
    int     tid;                /* thread ID */
    uint64  sp;                 /* saved stack pointer */
    uint64  pc;                 /* saved instruction pointer */
    struct thread_context context; /* saved registers */
};
```

Some parts of the struct may seem unnecessary, such as the `tid` in my code, and the `sp` and `pc` are already saved in the `context` field of type `struct thread_context`, which has a similar definition to that of the `struct context` used for processes, but we just forgot to update the redundant parts and it does not affect the program's execution. For the `context`, in fact, we can simply copy the definition from `struct context` to `thread_context`.

The code for `thread_schedule` looks similar to the `scheduler` used for process scheduling, so we can modify it in a similar way. To switch between threads, we need to implement the `thread_switch` function, which is defined in `uthread_switch.S`. We can easily update the assembly code by copying the code from `swtch.S`, which handles context switching for processes. Since we already copied the structure of the `context` in the previous step, the implementation should be similar.

Finally, we just need to initialize a newly created thread, which involves two parts: initializing the thread's `stack` and initializing the thread's `context`. In this implementation, the thread's stack is contained within the thread's structure, which makes it easier to control the thread.

To initialize the thread's `context`, we set the stack pointer (`sp`) to the top of the stack and the program counter (`pc`) to the function `func` that will be executed by the thread.

```
t->sp = (uint64)&t->stack[STACK_SIZE];
t->pc = (uint64)func;
```

To complete the initialization of the thread's `context`, we need to initialize the `sp` and `pc` again (due to the initial redundant design) and the other callee-save registers (`s0-s11`). After doing this, the task is finished.

Using threads (moderate)

The tasks at hand are not labs for the xv6 operating system, but rather programs that run directly on the host machine.

Our goal is to improve the performance of filling and updating the hash table using multithreading in C. The original code works correctly when run with one thread, but there are issues when using multiple threads. Our task is to resolve these issues and optimize the process using multithreading techniques.

Question

The lab website poses a question about the problem. The answer to this question can be found below:

1. **Why are there missing keys with 2 threads, but not with 1 thread? Identify a sequence of events with 2 threads that can lead to a key being missing. Submit your sequence with a short explanation in answers-thread.txt**

If multiple threads execute the program, race conditions may occur due to the lack of locks on the hash table when threads access it simultaneously. This can lead to conflicting updates and unpredictable behavior, such as when multiple threads try to insert or update entries in the table at the same time.

For example, consider a scenario where `thread 0` and `thread 1` both try to add a key-value pair to `bucket 0`. It may be difficult to determine what happened during the insert process, which could potentially result in many entries being lost from the linked entry list.

For instance, suppose the hash table initially has the following structure:

```
index  pointers  entries
  0    [  --]--> [(k1, v1), --]--> [(k2, v2), --]--> NULL
  1    [  --]--> NULL
  ...
```

Now, let's say that `thread 0` tries to insert a key-value pair `(k3, v3)` into `bucket 0`, while `thread 1` also wants to insert a key-value pair `(k4, v4)` into the same bucket. If `thread 0` enters the insert function first and executes up to the line `e->next = n;` before the entry pointer is updated to link to the newly inserted entry struct, then `thread 0` may switch to `thread 1`, which also enters the insert function and successfully updates the entry pointer to point to the entry with key pair `(k4, v4)`.

As a result, the structure of the first bucket would look like this:

```

0    [  --]--> [(k4, v4), --]--> [(k2, v2), --]--> NULL
                                ^
                                |
                                |
                        [(k3, v3), --]-->

```

When the switch back to thread 0 and it continues executing the code `*p = e`, it could cause the entry `[(k4, v4), ----]---->` to be lost. This is the cause of missing keys in this program.

To solve this issue, we can add a lock to each bucket to prevent race conditions. This way, multiple threads can work concurrently on different buckets without causing missing entries. Different approaches can be taken to implement this solution, but the provided lock and unlock statements make it easy to implement this solution.

Implementation

As we have previously discussed, the solution to this issue is straightforward: we can add a lock to each hash bucket to prevent race conditions. The provided instructions include some functions from `pthread.h` that can manage locks, and we can refer to the manual to learn how to use them.

```

pthread_mutex_t lock;           // declare a lock
pthread_mutex_init(&lock, NULL); // initialize the lock
pthread_mutex_lock(&lock);      // acquire lock
pthread_mutex_unlock(&lock);    // release lock

```

To implement this solution, we need to initialize the `locks` array and do so in the `main` function. Remember to also destroy the locks at the end of the `main` function.

```

pthread_mutex_t locks[NBUCKET]; // This is a global mutex locks array

...
// other part of code
...

int
main(int argc, char *argv[])
{
    pthread_t *tha;
    void *value;
    double t1, t0;

    for (int i = 0; i < NBUCKET; i++) {
        pthread_mutex_init(&locks[i], NULL); // initialize each lock
    }

    ...
    // other part of code
    ...

    for (int i = 0; i < NBUCKET; i++) {
        pthread_mutex_destroy(&locks[i]); // destroy each lock
    }
}

```

Then, we can add a lock corresponding to the computed hash index by adding a lock after this line

`int i = key % NBUCKET;` in the `put` function, which is the point where race conditions may occur. The updated `put` function would look like this:

```
static
void put(int key, int value)
{
    int i = key % NBUCKET;

    pthread_mutex_lock(&locks[i]);
    // is the key already present?
    struct entry *e = 0;
    for (e = table[i]; e != 0; e = e->next) {
        if (e->key == key)
            break;
    }
    if(e){
        // update the existing key.
        e->value = value;
    } else {
        // the new is new.
        insert(key, value, &table[i], table[i]);
    }

    pthread_mutex_unlock(&locks[i]);
}
```

After making these modifications, the tests should pass.

Barrier(moderate)

This task also introduces some functions from `pthread.h` that can manage conditional waiting and broadcasting for threads.

```
pthread_cond_wait(&cond, &mutex); // go to sleep on cond, releasing lock mutex, acquiring upon wake up
pthread_cond_broadcast(&cond);    // wake up every thread sleeping on cond
```

We only need to implement the `barrier` function, which should allow all threads to wait at certain points and resume when a certain condition is satisfied.

```
static void
barrier()
{
    pthread_mutex_lock(&bstate.barrier_mutex);
    bstate.nthread++;
    while(bstate.nthread != nthread && bstate.nthread != 0){
        pthread_cond_wait(&bstate.barrier_cond, &bstate.barrier_mutex);
    }
    if (bstate.nthread == nthread) {
        pthread_cond_broadcast(&bstate.barrier_cond);
        bstate.nthread = 0;
        bstate.round++;
    }
    pthread_mutex_unlock(&bstate.barrier_mutex);
}
```

The code for the `barrier` function is straightforward. When a thread enters the `barrier` from outside, we add a lock and increase the number of threads recorded in the global `bstate.nthread`. There are two conditions that the waiting loop must satisfy:

1. If the number of threads reaches `nthread`, the waiting loop should be skipped and the first thread to reach the code line after the while loop should be responsible for broadcasting to the other threads and starting the next round. The `bstate.nthread` should also be reset to zero.
2. When other threads are awoken by the thread that skips the while loop, they should be able to leave the loop. However, if `bstate.nthread` has already been updated by the last thread, the other threads may not leave the loop if we don't add the condition `bstate.nthread != 0`.

However, the code above has a bug that can cause a deadlock when we run the program with multiple threads. To debug this issue, we can print out the execution sequence of a deadlock session and manually walk through it step by step as follows:

Analysis of a deadlock in the `barrier()` function.

The deadlock occurs when the last thread to reach the barrier is the last thread to exit the `while` loop. Terminal output:

```
===== Entering barrier() === // Thread 0 enters barrier
nthread = 0, round = 0         // Thread 0 acquires lock
Entering cond_wait:            // Thread 0 enters cond_wait
    nthread = 1, round = 0     // Thread 0 releases lock and waits
===== Entering barrier() === // Thread 1 enters barrier
nthread = 1, round = 0         // Thread 1 acquires lock
Exited while loop:            // Thread 1 skips cond_wait
    nthread = 2, round = 0
Broadcasting...                // Thread 1 broadcasts
Broadcasted!
===== Exiting barrier() ==== // Thread 1 exits barrier, releases lock
Exiting cond_wait:            // Thread 0 wakes up, acquires lock
    nthread = 0, round = 1
Exited while loop:            // Thread 0 satisfies while loop condition and exits
    nthread = 0, round = 1
===== Exiting barrier() ==== // Thread 0 exits barrier, releases lock
===== Entering barrier() === // Thread 0 enters barrier
nthread = 0, round = 1         // Thread 0 acquires lock
Entering cond_wait:            // Thread 0 enters cond_wait
    nthread = 1, round = 1     // Thread 0 releases lock and waits
===== Entering barrier() === // Thread 1 enters barrier
nthread = 1, round = 1         // Thread 1 acquires lock
Exited while loop:            // Thread 1 skips cond_wait
    nthread = 2, round = 1
Broadcasting...                // Thread 1 broadcasts
Broadcasted!
===== Exiting barrier() ==== // Thread 1 exits barrier, releases lock
===== Entering barrier() === // Thread 1 enters barrier
nthread = 0, round = 2         // Thread 1 acquires lock
Entering cond_wait:            // Thread 1 enters cond_wait
    nthread = 1, round = 2
Exiting cond_wait:            // Bugs here, if a thread broadcasted
                                // other threads and then raced in the next round
    nthread = 1, round = 2     // and reached the barrier before the last thread
                                // exited the while loop, the last thread
Entering cond_wait:            // would be stuck in the cond_wait.
    nthread = 1, round = 2
```

Upon further examination, we discover that there is still a possibility for deadlock, as mentioned in the instructions (which we didn't fully understand at first):

"You have to handle the case in which one thread races around the loop before the others have exited the barrier. In particular, you are re-using the `bstate.nthread` variable from one round to the next. Make sure that a thread that leaves the barrier and races around the loop doesn't increase `bstate.nthread` while a previous round is still using it. "

This means that if a thread broadcasts to the other threads and then races into the next round and reaches the `while` loop in the `barrier` function before the last thread exits the `while` loop, the last thread could get stuck in the `cond_wait`.

To fix this issue, we can simply add a local variable `current_round` to temporarily store the round number that the thread enters the `while` loop. When the thread is awoken by another thread, it should first check the round number and, if it has changed (indicating that the next round has started before the thread has exited the `while` loop), it should just `break` out of the loop.

After making this modification, all tests should pass.

Make grade

```
make[1]: Leaving directory '/home/lydia/projects/xv6-labs-2021'
== Test uthread ==
$ make qemu-gdb
uthread: OK (3.8s)
== Test answers-thread.txt == answers-thread.txt: OK
== Test ph_safe == make[1]: Entering directory '/home/lydia/projects/xv6-labs-2021'
gcc -o ph -g -O2 -DSOL_THREAD -DLAB_THREAD notxv6/ph.c -pthread
make[1]: Leaving directory '/home/lydia/projects/xv6-labs-2021'
ph_safe: OK (16.1s)
== Test ph_fast == make[1]: Entering directory '/home/lydia/projects/xv6-labs-2021'
make[1]: 'ph' is up to date.
make[1]: Leaving directory '/home/lydia/projects/xv6-labs-2021'
ph_fast: OK (36.0s)
== Test barrier == make[1]: Entering directory '/home/lydia/projects/xv6-labs-2021'
gcc -o barrier -g -O2 -DSOL_THREAD -DLAB_THREAD notxv6/barrier.c -pthread
make[1]: Leaving directory '/home/lydia/projects/xv6-labs-2021'
barrier: OK (2.7s)
== Test time ==
time: OK
Score: 60/60
lydia@ubuntu-22-hp-040f1b4d:~/projects/xv6-labs-2021$
```

Reference

1. Xv6 Book Chapter 7: Scheduling - <https://pdos.csail.mit.edu/6.828/2021/xv6/book-riscv-rev2.pdf>
2. MIT 6.S081 Lab: Multithreading - <https://pdos.csail.mit.edu/6.828/2021/labs/thread.html>
3. Integer Calling convention - <https://tclin914.github.io/77838749/>
4. Lock and Unlock A Mutex - https://pubs.opengroup.org/onlinepubs/007908799/xsh/pthread_mutex_lock.html
5. Initialise or Destroy A Mutex - https://pubs.opengroup.org/onlinepubs/007908799/xsh/pthread_mutex_init.html
6. Thread Creation - https://pubs.opengroup.org/onlinepubs/007908799/xsh/pthread_create.html

7. Wait on A Condition - https://pubs.opengroup.org/onlinepubs/007908799/xsh/pthread_cond_wait.html
8. Signal or Broadcast A Condition - https://pubs.opengroup.org/onlinepubs/007908799/xsh/pthread_cond_broadcast.html